

```
In [1]: import pandas as pd
import numpy as np
```

Task 1

```
In [2]: df = pd.read_csv('insurance.csv', index_col=None)
df
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
995	39	female	23.275	3	no	northeast	7986.47525
996	39	female	34.100	3	no	southwest	7418.52200
997	63	female	36.850	0	no	southeast	13887.96850
998	33	female	36.290	3	no	northeast	6551.75010
999	36	female	26.885	0	no	northwest	5267.81815

1000 rows x 7 columns

```
In [3]: df.info
```

```
Out[3]: <bound method DataFrame.info of
smoker      region      charges      age      sex      bmi      children
0         19    female    27.900         0     yes    southwest    16884.92400
1         18     male    33.770         1      no     southeast    1725.55230
2         28     male    33.000         3      no     southeast    4449.46200
3         33     male    22.705         0      no     northwest    21984.47061
4         32     male    28.880         0      no     northwest    3866.85520
..      ...      ...      ...      ...      ...      ...      ...
995        39    female    23.275         3      no     northeast    7986.47525
996        39    female    34.100         3      no     southwest    7418.52200
997        63    female    36.850         0      no     southeast    13887.96850
998        33    female    36.290         3      no     northeast    6551.75010
999        36    female    26.885         0      no     northwest    5267.81815
```

[1000 rows x 7 columns]>

In [4]: `from sklearn.preprocessing import LabelEncoder`

```
LE = LabelEncoder()
df['sex'] = LE.fit_transform(df['sex'])
df['smoker'] = LE.fit_transform(df['smoker'])
df['region'] = LE.fit_transform(df['region'])

df
```

Out [4]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520
...
995	39	0	23.275	3	0	0	7986.47525
996	39	0	34.100	3	0	3	7418.52200
997	63	0	36.850	0	0	2	13887.96850
998	33	0	36.290	3	0	0	6551.75010
999	36	0	26.885	0	0	1	5267.81815

1000 rows × 7 columns

Label Encoding Info:

sex: Male - 1, Female - 0

smoker: Yes - 1, No - 0

region: southwest - 3, southeast - 2, northwest - 1, northeast - 0

In [5]: `data = df.iloc[:, :6].values`
`data`

Out [5]: `array([[19. , 0. , 27.9 , 0. , 1. , 3.],`
 `[18. , 1. , 33.77, 1. , 0. , 2.],`
 `[28. , 1. , 33. , 3. , 0. , 2.],`
 `...,`
 `[63. , 0. , 36.85, 0. , 0. , 2.],`
 `[33. , 0. , 36.29, 3. , 0. , 0.],`
 `[36. , 0. , 26.885, 0. , 0. , 1.]])`

```
In [6]: charges = df.iloc[:, -1].values
charges
```

```
Out[6]: array([16884.924 , 1725.5523 , 4449.462 , 21984.47061 ,
        3866.8552 , 3756.6216 , 8240.5896 , 7281.5056 ,
        6406.4107 , 28923.13692 , 2721.3208 , 27808.7251 ,
        1826.843 , 11090.7178 , 39611.7577 , 1837.237 ,
        10797.3362 , 2395.17155 , 10602.385 , 36837.467 ,
        13228.84695 , 4149.736 , 1137.011 , 37701.8768 ,
        6203.90175 , 14001.1338 , 14451.83515 , 12268.63225 ,
        2775.19215 , 38711. , 35585.576 , 2198.18985 ,
        4687.797 , 13770.0979 , 51194.55914 , 1625.43375 ,
        15612.19335 , 2302.3 , 39774.2763 , 48173.361 ,
        3046.062 , 4949.7587 , 6272.4772 , 6313.759 ,
        6079.6715 , 20630.28351 , 3393.35635 , 3556.9223 ,
        12629.8967 , 38709.176 , 2211.13075 , 3579.8287 ,
        23568.272 , 37742.5757 , 8059.6791 , 47496.49445 ,
        13607.36875 , 34303.1672 , 23244.7902 , 5989.52365 ,
        8606.2174 , 4504.6624 , 30166.61817 , 4133.64165 ,
        14711.7438 , 1743.214 , 14235.072 , 6389.37785 ,
        5920.1041 , 17663.1442 , 16577.7795 , 6799.458 ,
        11741.726 , 11946.6259 , 7726.854 , 11356.6609 ,
        2047.4121 , 1522.4607 , 2755.02005 , 6571.02425 ,
```

```
In [7]: from sklearn import linear_model
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [8]: regr = linear_model.LinearRegression()

regr.fit(data, charges)

y_pred = regr.predict(data)
```

```
In [9]: print("Coefficients: \n", regr.coef_)
print("Mean squared error: %.2f" % mean_squared_error(charges, y_pr
print("Coefficient of determination: %.2f" % r2_score(charges, y_pr
```

```
Coefficients:
[ 265.00281538 -276.07122609  332.34748636  415.6631237
 23799.08282857 -462.08574687]
Mean squared error: 34925480.46
Coefficient of determination: 0.76
```

Task 2

```
In [10]: df = pd.read_csv('Fraud.csv', index_col=None)
df
```

Out[10]:

	step	type	amount	nameOrig	oldbalanceOrg	newbalanceOrig	nan
0	1	PAYMENT	9839.64	C1231006815	170136.00	160296.36	M1979
1	1	PAYMENT	1864.28	C1666544295	21249.00	19384.72	M2044
2	1	TRANSFER	181.00	C1305486145	181.00	0.00	C553
3	1	CASH_OUT	181.00	C840083671	181.00	0.00	C38
4	1	PAYMENT	11668.14	C2048537720	41554.00	29885.86	M1230
...
1047295	95	PAYMENT	9028.31	C351212668	105818.00	96789.69	M1938
1047296	95	PAYMENT	6290.51	C404585795	96789.69	90499.18	M391
1047297	95	PAYMENT	2499.92	C2079661953	90499.18	87999.25	M227
1047298	95	PAYMENT	4780.32	C2115046838	802.00	0.00	M1154
1047299	95	PAYMENT	789.77	C386614070	30791.00	30001.23	M919

1047300 rows × 11 columns

```
In [11]: df.info
```

```
Out[11]: <bound method DataFrame.info of
nameOrig oldbalanceOrg newbalanceOrig \
0 1 PAYMENT 9839.64 C1231006815 170136.00
160296.36
1 1 PAYMENT 1864.28 C1666544295 21249.00
19384.72
2 1 TRANSFER 181.00 C1305486145 181.00
0.00
3 1 CASH_OUT 181.00 C840083671 181.00
0.00
4 1 PAYMENT 11668.14 C2048537720 41554.00
29885.86
... ...
...
1047295 95 PAYMENT 9028.31 C351212668 105818.00
96789.69
1047296 95 PAYMENT 6290.51 C404585795 96789.69
90499.18
1047297 95 PAYMENT 2499.92 C2079661953 90499.18
87999.25
1047298 95 PAYMENT 4780.32 C2115046838 802.00
0.00
1047299 95 PAYMENT 789.77 C386614070 30791.00
30001.23
```

```
laggedFraud
0      M1979787155      0.0      0.0      0
0
1      M2044282225      0.0      0.0      0
0
2      C553264065      0.0      0.0      1
0
3      C38997010      21182.0      0.0      1
0
4      M1230701703      0.0      0.0      0
0
...      ...      ...      ...      ...
...
1047295 M1938858197      0.0      0.0      0
0
1047296 M391493091      0.0      0.0      0
0
1047297 M227905058      0.0      0.0      0
0
1047298 M1154565434      0.0      0.0      0
0
1047299 M919485291      0.0      0.0      0
0

[1047300 rows x 11 columns]>
```

In [12]: `from sklearn.preprocessing import LabelEncoder`

```
LE = LabelEncoder()
df['type'] = LE.fit_transform(df['type'])
df['nameOrig'] = LE.fit_transform(df['nameOrig'])
df['nameDest'] = LE.fit_transform(df['nameDest'])

df
```

Out [12]:

	step	type	amount	nameOrig	oldbalanceOrg	newbalanceOrig	nameDest	oldbalanceDest
0	1	3	9839.64	125004	170136.00	160296.36	274496	
1	1	3	1864.28	360882	21249.00	19384.72	286268	
2	1	4	181.00	165023	181.00	0.00	73503	
3	1	1	181.00	960498	181.00	0.00	65423	
4	1	3	11668.14	567221	41554.00	29885.86	137860	
...
1047295	95	3	9028.31	694724	105818.00	96789.69	266884	
1047296	95	3	6290.51	723239	96789.69	90499.18	337296	
1047297	95	3	2499.92	583942	90499.18	87999.25	307566	
1047298	95	3	4780.32	603263	802.00	0.00	124109	
1047299	95	3	789.77	713704	30791.00	30001.23	433968	

1047300 rows × 11 columns

In [13]: `data = df.iloc[:, :10].values`
`data`

Out [13]: `array([[1.00000e+00, 3.00000e+00, 9.83964e+03, ..., 0.00000e+00, 0.00000e+00, 0.00000e+00],`
`[1.00000e+00, 3.00000e+00, 1.86428e+03, ..., 0.00000e+00, 0.00000e+00, 0.00000e+00],`
`[1.00000e+00, 4.00000e+00, 1.81000e+02, ..., 0.00000e+00, 0.00000e+00, 1.00000e+00],`
`...,`
`[9.50000e+01, 3.00000e+00, 2.49992e+03, ..., 0.00000e+00, 0.00000e+00, 0.00000e+00],`
`[9.50000e+01, 3.00000e+00, 4.78032e+03, ..., 0.00000e+00, 0.00000e+00, 0.00000e+00],`
`[9.50000e+01, 3.00000e+00, 7.89770e+02, ..., 0.00000e+00, 0.00000e+00, 0.00000e+00]])`

```
In [14]: label = df.iloc[:, -2].values
label
```

```
Out[14]: array([0, 0, 1, ..., 0, 0, 0])
```

```
In [15]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(data, label, te
```

```
In [16]: from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression(random_state=16)
logreg.fit(X_train, y_train)
```

```
Out[16]: LogisticRegression(random_state=16)
```

```
In [17]: y_pred = logreg.predict(X_test)
```

```
In [18]: from sklearn import metrics
cm = metrics.confusion_matrix(y_test, y_pred)
cm
```

```
Out[18]: array([[209217,    22],
               [   138,    83]])
```

```
In [19]: from sklearn.metrics import classification_report
target_names = ['not spam', 'spam']
print(classification_report(y_test, y_pred, target_names=target_nam
```

	precision	recall	f1-score	support
not spam	1.00	1.00	1.00	209239
spam	0.79	0.38	0.51	221
accuracy			1.00	209460
macro avg	0.89	0.69	0.75	209460
weighted avg	1.00	1.00	1.00	209460

```
In [20]: acc = metrics.accuracy_score(y_test, y_pred)
print("Accuracy: "+str(acc))
```

```
Accuracy: 0.9992361310035329
```

